

# BitePulse AI - Modeling: MS-TCN with Frame-Level Supervision

- Name/Group: Aktham Almomani / Group 13

## Introduction

Instead of classifying short video windows with a heavy 3D-CNN, we move to an MS-TCN (Multi-Stage Temporal Convolutional Network) that operates on **frame-level pose features**. The original dataset has **16 action labels**, but for BitePulse we collapse them into a binary target: **INTAKE** (frames labeled *eat it*) vs **NON-INTAKE** (all other labels). The MS-TCN takes the full sequence of frame features and predicts this binary label at each time step, giving us a lightweight, sequence-aware model that we can later convert into bite events and pacing metrics.

## Imports and basic setup

```
# Required imports:  
import collections  
import itertools  
import os  
from pathlib import Path  
import shutil  
import sys  
import time  
from collections import Counter  
import math  
import copy  
  
import matplotlib.pyplot as plt  
import numpy as np  
import pandas as pd  
import seaborn as sns  
from sklearn.model_selection import StratifiedShuffleSplit,  
train_test_split  
from sklearn.metrics import (  
    accuracy_score,  
    auc,  
    classification_report,  
    confusion_matrix,  
    f1_score,  
    precision_recall_curve,  
    precision_recall_fscore_support,  
    recall_score,  
    roc_curve,  
)
```

```

from sklearn.model_selection import KFold
from sklearn.preprocessing import label_binarize
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.utils.data import DataLoader, Dataset

pd.set_option('display.max_columns', None)
pd.set_option('display.width', 1000)

# Mount Google Drive:
from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

```

## Starting Fresh - Copy Original EatSense Dataset

```

# Set paths:
SRC = Path("/content/drive/MyDrive/eatsense")
DST = Path("/content/drive/MyDrive/eatsense_v1")

# Folders that we need - original EatSense data:
raw_dirs = [
    "rgb",
    "depth",
    "poses_true",
    "poses_fake",
    "true2d_parquet",
    "labels_v0",
    "all_misc",
    "test",
]
# Extras just in case:
raw_files = [
    "manifest_eatsense_rgb_true2d.csv",
    "segments_index_v0.parquet",
    "frames_index_v0.parquet",
    "splits_v0.csv",
    "fps_by_key.json",
]

```

```

# Copy folders:
for d in raw_dirs:
    src_d = SRC / d
    dst_d = DST / d
    if src_d.exists():
        print(f"Copying dir: {src_d} -> {dst_d}")
        shutil.copytree(src_d, dst_d)
    else:
        print(f"SKIP (dir not found): {src_d}")

# Copy files
for f in raw_files:
    src_f = SRC / f
    dst_f = DST / f
    if src_f.exists():
        print(f"Copying file: {src_f} -> {dst_f}")
        shutil.copy2(src_f, dst_f)
    else:
        print(f"SKIP (file not found): {src_f}")

print("\nDone. Clean raw copy is at:", DST)

Copying dir: /content/drive/MyDrive/eatsense/rgb ->
/content/drive/MyDrive/eatsense_v1/rgb
Copying dir: /content/drive/MyDrive/eatsense/depth ->
/content/drive/MyDrive/eatsense_v1/depth
Copying dir: /content/drive/MyDrive/eatsense/poses_true ->
/content/drive/MyDrive/eatsense_v1/poses_true
Copying dir: /content/drive/MyDrive/eatsense/poses_fake ->
/content/drive/MyDrive/eatsense_v1/poses_fake
Copying dir: /content/drive/MyDrive/eatsense/true2d_parquet ->
/content/drive/MyDrive/eatsense_v1/true2d_parquet
Copying dir: /content/drive/MyDrive/eatsense/labels_v0 ->
/content/drive/MyDrive/eatsense_v1/labels_v0
Copying dir: /content/drive/MyDrive/eatsense/all_misc ->
/content/drive/MyDrive/eatsense_v1/all_misc
Copying dir: /content/drive/MyDrive/eatsense/test ->
/content/drive/MyDrive/eatsense_v1/test
Copying file:
/content/drive/MyDrive/eatsense/manifest_eatsense_rgb_true2d.csv ->
/content/drive/MyDrive/eatsense_v1/manifest_eatsense_rgb_true2d.csv
Copying file:
/content/drive/MyDrive/eatsense/segments_index_v0.parquet ->
/content/drive/MyDrive/eatsense_v1/segments_index_v0.parquet
Copying file: /content/drive/MyDrive/eatsense/frames_index_v0.parquet
-> /content/drive/MyDrive/eatsense_v1/frames_index_v0.parquet
Copying file: /content/drive/MyDrive/eatsense/splits_v0.csv ->
/content/drive/MyDrive/eatsense_v1/splits_v0.csv
Copying file: /content/drive/MyDrive/eatsense/fps_by_key.json ->
/content/drive/MyDrive/eatsense_v1/fps_by_key.json

```

```

Done. Clean raw copy is at: /content/drive/MyDrive/eatsense_v1
ROOT = Path("/content/drive/MyDrive/eatsense_v1")

for p in ROOT.iterdir():
    print("[DIR] " if p.is_dir() else "[FILE]" + " ", p.name)

[DIR]    rgb
[DIR]    depth
[DIR]    poses_true
[DIR]    poses_fake
[DIR]    true2d_parquet
[DIR]    labels_v0
[DIR]    all_misc
[DIR]    test
[FILE]   manifest_eatsense_rgb_true2d.csv
[FILE]   segments_index_v0.parquet
[FILE]   frames_index_v0.parquet
[FILE]   splits_v0.csv
[FILE]   fps_by_key.json

```

## Environment Setup and Data Schema Inspection

```

# our roots:
base_path = '/content/drive/MyDrive/eatsense_v1'

if os.path.exists(base_path):
    print(f">>> Dataset root found at: {base_path}")
else:
    print(f">>> Dataset root NOT found at: {base_path}")
    sys.exit(1)

>>> Dataset root found at: /content/drive/MyDrive/eatsense_v1

def locate_true_3d_files(base_path):
    # Constructing the path based on the raw structure tree
    true_3d_dir = os.path.join(base_path, 'poses_true', 'True3D')

    print(f"\n>>> [NAV] Locating True3D directory: {true_3d_dir}")

    if not os.path.exists(true_3d_dir):
        print("=>> 'poses_true/True3D' directory does not exist.")
        sys.exit(1)

    csv_files = sorted([f for f in os.listdir(true_3d_dir) if
f.endswith('.csv')])

    if not csv_files:
        print("=>> No CSV files found in True3D directory.")
        sys.exit(1)

```

```

print(f">>> Found {len(csv_files)} pose files.")
print("    Sample filenames:")
for f in csv_files[:3]:
    print(f"      - {f}")

return true_3d_dir, csv_files

def inspect_schema(file_dir, filename):
    full_path = os.path.join(file_dir, filename)
    print(f"\n>>> Loading file: {filename}")

    try:
        df = pd.read_csv(full_path)

        print("\n" + "*"*60)
        print(f"DATA INSPECTION REPORT: {filename}")
        print("*"*60)

        # 1. Dimensions:
        print(f"\n[1] Dimensions:")
        print(f"      Rows (Frames): {df.shape}")
        print(f"      Columns (Features): {df.shape[1]}")

        # 2. Column Headers:
        print(f"\n[2] Column Headers:")
        print(f"      {df.columns.tolist()}")

        # 3. Data Types:
        print(f"\n[3] Data Types Summary:")
        print(df.dtypes.value_counts())

        # 4. Null Value Check (Critical for back-projection validity):
        null_counts = df.isnull().sum()
        total_nulls = null_counts.sum()
        print(f"\n[4] Missing Data Check:")
        if total_nulls == 0:
            print("      No missing values detected.")
        else:
            print(f"      WARNING: {total_nulls} missing values
detected.")
            print(null_counts=null_counts > 0))

        # 5. Label Identification Strategy:
        # Searching for columns that might contain the 'Action'
labels:
    label_keywords = ['label', 'action', 'class', 'annotation',
'ground_truth']
    potential_labels = [c for c in df.columns if any(k in
c.lower() for k in label_keywords)]

```

```

        print(f"\n[5] Label Column Identification:")
        if potential_labels:
            print(f"    Potential label columns found:
{potential_labels}")
            for col in potential_labels:
                unique_vals = df[col].unique()
                print(f"        Unique values in '{col}'"
({len(unique_vals)}): {unique_vals[:10]}...")
            else:
                print("        WARNING: No explicit label column found based
on keywords.")
                print("        Check if labels are in a separate file or
encoded in the filename.")

        # 6. Data Sample:
        print(f"\n[6] Head (First 5 Rows):")
        print(df.head())
        print("=*60")

    except Exception as e:
        print(f">>> Could not read CSV: {e}")

if __name__ == "__main__":
    dataset_root = '/content/drive/MyDrive/eatsense_v1'

    # Locate Target Files:
    target_dir, file_list = locate_true_3d_files(dataset_root)

    # Testing the schema:
    test_filename = file_list[0]

    print(f">>> Selected test file: {test_filename}")

    # Run Inspection on that SINGLE file:
    inspect_schema(target_dir, test_filename)

>>> [NAV] Locating True3D directory:
/content/drive/MyDrive/eatsense_v1/poses_true/True3D
>>> Found 135 pose files.
Sample filenames:
- 20210518_230219.csv
- 20210523_202300.csv
- 20210529_150552.csv

>>> Selected test file: 20210518_230219.csv

>>> Loading file: 20210518_230219.csv
=====

```

DATA INSPECTION REPORT: 20210518\_230219.csv

[1] Dimensions:

Rows (Frames): (3668, 28)  
Columns (Features): 28

## [2] Column Headers:

```
[ 'Unnamed: 0', 'Project', 'Action', 'Img', 'head_x', 'head_y',  
'head_z', 'Mid-Shoulder_x', 'Mid-Shoulder_y', 'Mid-Shoulder_z',  
'Right-Shoulder_x', 'Right-Shoulder_y', 'Right-Shoulder_z', 'Right-  
Elbow_x', 'Right-Elbow_y', 'Right-Elbow_z', 'Right-Wrist_x', 'Right-  
Wrist_y', 'Right-Wrist_z', 'Left-shoulder_x', 'Left-shoulder_y',  
'Left-shoulder_z', 'Left-Elbow_x', 'Left-Elbow_y', 'Left-Elbow_z',  
'Left-Wrist_x', 'Left-Wrist_y', 'Left-Wrist_z' ]
```

### [3] Data Types Summary:

float64 24

int64 2

object 2

### Missing Data Check:

## [5] Label Column Identification

### Label Column Identification:

```
Potential label columns found: ['Action']
Unique values in 'Action' (7): ['other' 'pick up tools with both
hands']
```

'pick food from utensil with tools in both hands'

move hand towards mouth! Eat it! move hand away from mouth!

'move hand towards mouth  
'food in hand at table'']

[6] Head (First 5 Rows):

```
[6] head (first 5 rows):
   Unnamed: 0      Project Action  Img  head_x  head_y
head_z  Mid-Shoulder_x  Mid-Shoulder_y  Mid-Shoulder_z  Right-
Shoulder_x  Right-Shoulder_y  Right-Shoulder_z  Right-Elbow_x  Right-
Elbow_y  Right-Elbow_z  Right-Wrist_x  Right-Wrist_y  Right-Wrist_z
Left-shoulder_x  Left-shoulder_y  Left-shoulder_z  Left-Elbow_x  Left-
Elbow_y  Left-Elbow_z  Left-Wrist_x  Left-Wrist_y  Left-Wrist_z
0          0  20210518_230219  other      0  0.672821  0.470648
2.015      0.646179        0.452282           2.178
0.802871      0.440989        2.1920        0.886018
0.243377      2.3770        0.891254        0.024731        2.4550
0.515521      0.486317        2.2540        0.463108        0.315621
2.460      0.428864        0.061444        2.6680
1          1  20210518_230219  other      1  0.649279  0.476431
2.054      0.612705        0.452670           2.197
0.751921      0.441700        2.1780        0.841623
0.251605      2.3820        0.851481        0.028976        2.4780
```

0.478349	0.478318	2.2680	0.414328	0.276436
2.502	0.380771	0.043598	2.6340	
2	2 20210518_230219	other	2 0.593581	0.468498
2.034	0.552275	0.449785	2.183	
0.690729	0.444410		2.1740	0.787419
0.256642	2.3930	0.799468	0.032829	2.4660
0.428846	0.469359	2.2780	0.362557	0.264914
2.508	0.339043	0.009495	2.6410	
3	3 20210518_230219	other	3 0.542065	0.466655
2.030	0.499283	0.442714	2.183	
0.640596	0.446250	2.1805	0.723637	
0.255999	2.3845	0.737280	0.036315	2.4605
0.363528	0.454071	2.2610	0.319674	0.252407
2.505	0.303771	-0.017068	2.6545	
4	4 20210518_230219	other	4 0.500427	0.464289
2.032	0.457094	0.446365	2.190	
0.589766	0.449583	2.1760	0.675111	
0.256642	2.3900	0.708495	0.040635	2.4605
0.314445	0.440421	2.2535	0.275995	0.243434
2.505	0.262664	-0.007389	2.6375	
<hr/>				

## Define Input and Output Paths

First, let's define the base path for the dataset, the input directory for `poses_true/True3D` CSVs, and create output directories for the processed feature (`.npy`) and label (`.txt`) files.

```
# Our Dataset root:
dataset_root = Path('/content/drive/MyDrive/eatsense_v1')

# Let's build the input directory path for poses_true/True3D CSVs:
true_3d_input_dir = dataset_root / 'poses_true' / 'True3D'

# Here, let's define the output directory path for processed feature files:
features_output_dir = dataset_root / 'processed_data' / 'features'

# Then let's define the output directory path for processed label files:
labels_output_dir = dataset_root / 'processed_data' / 'labels'

# Finally, let's create the output directories if they don't exist:
features_output_dir.mkdir(parents=True, exist_ok=True)
labels_output_dir.mkdir(parents=True, exist_ok=True)

print(f"Dataset Root: {dataset_root}")
print(f"True3D Input Directory: {true_3d_input_dir}")
print(f"Features Output Directory: {features_output_dir}")
```

```

print(f"Labels Output Directory: {labels_output_dir}")
print("Output directories created/verified.")

Dataset Root: /content/drive/MyDrive/eatsense_v1
True3D Input Directory:
/content/drive/MyDrive/eatsense_v1/poses_true/True3D
Features Output Directory:
/content/drive/MyDrive/eatsense_v1/processed_data/features
Labels Output Directory:
/content/drive/MyDrive/eatsense_v1/processed_data/labels
Output directories created/verified.

```

## Process and Convert True3D Data

Alright, here let's iterate through each CSV file in the `poses_true/True3D` directory. For each file, let's load it, extract the coordinate columns as features, and the 'Action' column as labels. Convert the features to an array and save it as a `.npy` file. Save the labels as a plain text file, with each label on a new line.

```

# Alright, let's list of all CSV files in the true_3d_input_dir:
csv_files = sorted([f for f in os.listdir(true_3d_input_dir) if f.endswith('.csv')])

processed_count = 0
print(f"\n>>> Starting processing of {len(csv_files)} CSV files...")

for csv_file in csv_files:
    full_csv_path = true_3d_input_dir / csv_file
    print(f"    Processing: {csv_file}")

    # Load the CSV file into df:
    df = pd.read_csv(full_csv_path)

    # Define columns to exclude from features:
    exclude_cols = ['Unnamed: 0', 'Project', 'Action', 'Img']
    feature_cols = [col for col in df.columns if col not in exclude_cols]

    # Extract coordinate columns as features and convert to array:
    features = df[feature_cols].to_numpy()

    # Extract 'Action' column as labels:
    labels = df['Action'].tolist()

    # Define output filenames:
    base_filename = csv_file.replace('.csv', '')
    feature_output_filename = f"{base_filename}.npy"
    label_output_filename = f"{base_filename}.txt"

```

```
# Save features as .npy file:
np.save(features_output_dir / feature_output_filename, features)

# Save labels as a text file (one label per line):
with open(labels_output_dir / label_output_filename, 'w') as f:
    for label in labels:
        f.write(f"{label}\n")

processed_count += 1

print(f"\n>>> Successfully processed {processed_count} files.")
print(f"Features saved to: {features_output_dir}")
print(f"Labels saved to: {labels_output_dir}")
```

```
>>> Starting processing of 135 CSV files...
Processing: 20210518_230219.csv
Processing: 20210523_202300.csv
Processing: 20210529_150552.csv
Processing: 20210529_153708.csv
Processing: 20210530_153343.csv
Processing: 20210531_150448.csv
Processing: 20210603_130948.csv
Processing: 20210605_155355.csv
Processing: 20210606_154234.csv
Processing: 20210607_152559.csv
Processing: 20210607_154140.csv
Processing: 20210608_163819.csv
Processing: 20210609_154241.csv
Processing: 20210609_215756.csv
Processing: 20210609_220502.csv
Processing: 20210609_221133.csv
Processing: 20210610_144155.csv
Processing: 20210610_230110.csv
Processing: 20210616_163847.csv
Processing: 20210616_170258.csv
Processing: 20210617_165755.csv
Processing: 20210617_224119.csv
Processing: 20210620_150313.csv
Processing: 20210620_230306.csv
Processing: 20210620_230707.csv
Processing: 20210621_125816.csv
Processing: 20210627_142532.csv
Processing: 20210714_125156.csv
Processing: 20210716_151219.csv
Processing: 20210719_135251.csv
Processing: 20210728_144106.csv
Processing: 20210804_144735.csv
Processing: 20210804_145925.csv
Processing: 20210810_163907.csv
```

Processing: 20210811\_130812.csv  
Processing: 20210811\_132135.csv  
Processing: 20210813\_195607.csv  
Processing: 20210814\_210449.csv  
Processing: 20210816\_191136.csv  
Processing: 20210923\_155848.csv  
Processing: 20210923\_161107.csv  
Processing: 20210923\_163035.csv  
Processing: 20210923\_163801.csv  
Processing: 20211001\_152746.csv  
Processing: 20211001\_154310.csv  
Processing: 20211029\_144948.csv  
Processing: 20211029\_150526.csv  
Processing: 20220809\_125236.csv  
Processing: 20220809\_142839.csv  
Processing: 20220809\_144052.csv  
Processing: 20220811\_113712.csv  
Processing: 20220811\_114634.csv  
Processing: 20220811\_121550.csv  
Processing: 20220811\_122341.csv  
Processing: 20220811\_124226.csv  
Processing: 20220812\_112034.csv  
Processing: 20220812\_113020.csv  
Processing: 20220812\_115050.csv  
Processing: 20220812\_120241.csv  
Processing: 20220812\_121833.csv  
Processing: 20220812\_122535.csv  
Processing: 20220812\_123723.csv  
Processing: 20220812\_124345.csv  
Processing: 20220812\_125132.csv  
Processing: 20220812\_141747.csv  
Processing: 20221013\_124150.csv  
Processing: 20221013\_125118.csv  
Processing: 20221013\_131254.csv  
Processing: 20221013\_132443.csv  
Processing: 20221013\_134216.csv  
Processing: 20221013\_135010.csv  
Processing: 20221013\_142709.csv  
Processing: 20221013\_143133.csv  
Processing: 20221014\_113751.csv  
Processing: 20221014\_114952.csv  
Processing: 20221014\_120757.csv  
Processing: 20221014\_122331.csv  
Processing: 20221014\_124303.csv  
Processing: 20221014\_125502.csv  
Processing: 20221014\_140942.csv  
Processing: 20221014\_141711.csv  
Processing: 20221014\_143649.csv  
Processing: 20221014\_144233.csv

Processing: 20230109\_114530.csv  
Processing: 20230109\_115040.csv  
Processing: 20230109\_115716.csv  
Processing: 20230109\_120255.csv  
Processing: 20230109\_123622.csv  
Processing: 20230109\_124022.csv  
Processing: 20230109\_124729.csv  
Processing: 20230109\_125206.csv  
Processing: 20230109\_132632.csv  
Processing: 20230109\_133112.csv  
Processing: 20230109\_133609.csv  
Processing: 20230109\_134123.csv  
Processing: 20230110\_113650.csv  
Processing: 20230110\_113938.csv  
Processing: 20230110\_114318.csv  
Processing: 20230110\_114618.csv  
Processing: 20230110\_122933.csv  
Processing: 20230110\_123332.csv  
Processing: 20230110\_123733.csv  
Processing: 20230110\_124143.csv  
Processing: 20230110\_133309.csv  
Processing: 20230110\_133944.csv  
Processing: 20230110\_134822.csv  
Processing: 20230110\_135409.csv  
Processing: 20230111\_114057.csv  
Processing: 20230111\_114640.csv  
Processing: 20230111\_115239.csv  
Processing: 20230111\_115909.csv  
Processing: 20230124\_121223.csv  
Processing: 20230124\_121646.csv  
Processing: 20230124\_122123.csv  
Processing: 20230124\_122705.csv  
Processing: 20230124\_130607.csv  
Processing: 20230124\_131109.csv  
Processing: 20230124\_131740.csv  
Processing: 20230124\_132312.csv  
Processing: 20230124\_135539.csv  
Processing: 20230124\_135936.csv  
Processing: 20230124\_140432.csv  
Processing: 20230124\_141139.csv  
Processing: 20230125\_140902.csv  
Processing: 20230125\_141519.csv  
Processing: 20230125\_142103.csv  
Processing: 20230125\_142618.csv  
Processing: 20230127\_120259.csv  
Processing: 20230127\_120640.csv  
Processing: 20230127\_121212.csv  
Processing: 20230127\_121834.csv  
Processing: 20230131\_130800.csv

```

Processing: 20230131_131313.csv
Processing: 20230131_131949.csv
Processing: 20230131_132739.csv

>>> Successfully processed 135 files.
Features saved to:
/content/drive/MyDrive/eatsense_v1/processed_data/features
Labels saved to:
/content/drive/MyDrive/eatsense_v1/processed_data/labels

```

## Verify Converted Files

Now, let's verify by listing the contents of the newly created feature (.npy) and label (.txt) directories to confirm that the conversion process was successful and the files are correctly generated.

```

# List contents of the features_output_dir:
print(f"\n>>> Contents of {features_output_dir}:")
feature_files = sorted(os.listdir(features_output_dir))
if feature_files:
    print(f"Found {len(feature_files)} feature files. Sample
filenames:")
    for f in feature_files[:5]:
        print(f"    - {f}")
else:
    print("No feature files found.")

# List contents of the labels_output_dir:
print(f"\n>>> Contents of {labels_output_dir}:")
label_files = sorted(os.listdir(labels_output_dir))
if label_files:
    print(f"    Found {len(label_files)} label files. Sample
filenames:")
    for f in label_files[:5]:
        print(f"        - {f}")
else:
    print("No label files found.")

# load and inspect one .npy file and one .txt file:
if feature_files and label_files:
    sample_feature_file = features_output_dir / feature_files[0]
    sample_label_file = labels_output_dir / label_files[0]

    print(f"\n>>> Inspecting sample feature file:
{sample_feature_file.name}")
    loaded_features = np.load(sample_feature_file)
    print(f"Shape: {loaded_features.shape}")
    print(f"First 5 rows:\n{loaded_features[:5]}")

```

```
    print(f"\n>>> Inspecting sample label file:\n{sample_label_file.name}")
        with open(sample_label_file, 'r') as f:
            loaded_labels = [line.strip() for line in f.readlines()]
        print(f"Total labels: {len(loaded_labels)}")
        print(f"First 10 labels: {loaded_labels[:10]}\n")
else:
    print("\nSkipping sample file inspection as no files were found.\n")

>>> Contents of
/content/drive/MyDrive/eatsense_v1/processed_data/features:
Found 135 feature files. Sample filenames:
- 20210518_230219.npy
- 20210523_202300.npy
- 20210529_150552.npy
- 20210529_153708.npy
- 20210530_153343.npy

>>> Contents of
/content/drive/MyDrive/eatsense_v1/processed_data/labels:
Found 135 label files. Sample filenames:
- 20210518_230219.txt
- 20210523_202300.txt
- 20210529_150552.txt
- 20210529_153708.txt
- 20210530_153343.txt

>>> Inspecting sample feature file: 20210518_230219.npy
Shape: (3668, 24)
First 5 rows:
[[ 0.67282104  0.4706485   2.0150001   0.64617854  0.45228249
2.17800021
  0.80287057  0.44098946  2.19200015  0.88601834  0.24337713
2.37700009
  0.89125431  0.0247308   2.45500016  0.51552123  0.48631701
2.25400019
  0.46310815  0.31562111  2.46000004  0.4288637   0.0614443
2.66800022]
 [ 0.64927858  0.47643125  2.05400014  0.61270493  0.45266983
2.19700003
  0.75192106  0.44170031  2.17800021  0.84162337  0.25160468
2.38200021
  0.85148114  0.02897575  2.47800016  0.47834876  0.47831813
2.26800013
  0.41432756  0.27643633  2.50200009  0.38077128  0.04359762
2.63400006]
 [ 0.59358084  0.46849802  2.03400016  0.55227512  0.44978529
2.18300009
  0.69072932  0.44441     2.17400002  0.7874195   0.25664216
```

```

2.39300013
 0.79946768  0.03282926  2.46600008  0.42884567  0.4693591
2.27800012
 0.3625567   0.26491374  2.50800014  0.33904272  0.0094955
2.64100003]
 [ 0.54206467  0.46665531  2.03000009  0.49928251  0.44271427
2.18300009
 0.64059615  0.44624978  2.18050015  0.7236374   0.25599867
2.38450015
 0.73727953  0.03631539  2.46050012  0.36352819  0.45407113
2.26100016
 0.31967399  0.25240687  2.50500011  0.30377084 -0.01706792
2.65450013]
 [ 0.50042689  0.46428898  2.03200018  0.45709372  0.44636473
2.19000006
 0.58976591  0.4495829   2.17600012  0.67511129  0.25664216
2.3900001
 0.7084949   0.04063484  2.46050012  0.31444457  0.44042143
2.2535001
 0.2759949   0.24343412  2.50500011  0.26266438 -0.00738864
2.63750005]]

```

```

>>> Inspecting sample label file: 20210518_230219.txt
Total labels: 3668
First 10 labels: ['other', 'other', 'other', 'other', 'other',
'other', 'other', 'other', 'other', 'other']

```

## MS-TCN model: Frame-Level

Now we'll build an MS-TCN model **from scratch** for frame-level intake detection. The plan:

- 1. Prepare frame-level labels**

Build a `FrameDataset` and `collate_fn` that load pose features, create one label per frame, and pad variable-length sequences (using `-100` as the ignore index).

- 2. Define the MS-TCN architecture**

Implement a multi-stage temporal CNN that takes frame sequences as input and outputs a class logit for every frame (INTAKE vs NON\_INTAKE).

- 3. Handle class imbalance**

Compute class frequencies from the training set and use a **class-weighted `nn.CrossEntropyLoss`** so rare INTAKE frames get higher weight.

- 4. Write the training loop**

Train MS-TCN on batches of sequences: forward pass → frame-level loss → backprop → optimizer step, with validation at the end of each epoch.

## 5. Add evaluation utilities

Flatten frame-level predictions and labels to compute accuracy, precision, recall, F1, ROC AUC, PR AUC, and a confusion matrix.

## 6. Visualize performance

Plot **frame-level confusion matrix**, **ROC curve**, and **Precision-Recall curve** and interpret where the model does well or struggles.

## 7. Document the changes

Summarize the full MS-TCN pipeline (data → model → training → metrics) and how it improves over the earlier window-based TCN setup.

## Modify Dataset for Frame-Level Labels

Here, let's update the MSTCNDataset to ensure it correctly loads and prepares frame-level labels (one label per frame) for each sequence, rather than aggregating them for sequence-level classification.

```
# Re-load all unique labels:
all_unique_labels = set()
label_files_for_mapping = sorted(os.listdir(labels_output_dir))
for label_file in label_files_for_mapping:
    with open(labels_output_dir / label_file, 'r') as f:
        for line in f:
            all_unique_labels.add(line.strip())

sorted_unique_labels = sorted(list(all_unique_labels))
print(f"Found {len(sorted_unique_labels)} raw labels:
{sorted_unique_labels}")

# Define binary mapping: NON_INTAKE = 0, INTAKE = 1
INTAKE_LABELS = {"eat it"}

label_to_id = {
    "NON_INTAKE": 0,
    "INTAKE": 1
}
id_to_label = {v: k for k, v in label_to_id.items()}

print("Binary label_to_id:", label_to_id)
print("Binary id_to_label:", id_to_label)

# Custom Dataset that binarizes frame-level labels:
class MSTCNDataset(Dataset):
    def __init__(self, features_dir, labels_dir, label_to_id,
intake_labels):
        self.features_dir = features_dir
        self.labels_dir = labels_dir
        self.label_to_id = label_to_id
```

```

    self.intake_labels = intake_labels

    # Get sorted list of all file names:
    self.filenames = sorted([f.stem for f in
Path(features_dir).glob('*.*np'))])

    def __len__(self):
        return len(self.filenames)

    def __getitem__(self, idx):
        filename = self.filenames[idx]

        # Load features:
        features_path = self.features_dir / f"{filename}.npy"
        features = np.load(features_path)
        features_tensor = torch.from_numpy(features).float()      # (T,
F)
        features_tensor = features_tensor.transpose(0, 1)          # (F,
T) for Conv1d

        # Load original string labels:
        labels_path = self.labels_dir / f"{filename}.txt"
        with open(labels_path, 'r') as f:
            raw_labels = [line.strip() for line in f]

        # Map each raw label to binary 0/1:
        numerical_labels = [
            self.label_to_id["INTAKE"] if lbl in self.intake_labels
            else self.label_to_id["NON_INTAKE"]
            for lbl in raw_labels
        ]

        labels_tensor = torch.tensor(numerical_labels,
dtype=torch.long)  # (T,)
        return features_tensor, labels_tensor

def collate_fn(batch):
    max_seq_len = max([item[0].shape[1] for item in batch])

    padded_features = []
    padded_labels = []

    for features, labels in batch:
        pad_amount_features = max_seq_len - features.shape[1]
        padded_features.append(F.pad(features, (0,
pad_amount_features), 'constant', 0))

        pad_amount_labels = max_seq_len - labels.shape[0]
        padded_labels.append(F.pad(labels, (0, pad_amount_labels),

```

```

'constant', -100))

batch_features = torch.stack(padded_features)
batch_labels = torch.stack(padded_labels)

return batch_features, batch_labels

# Instantiate dataset & dataloader:
mstcn_dataset = MSTCNDataset(
    features_output_dir,
    labels_output_dir,
    label_to_id=label_to_id,
    intake_labels=INTAKE_LABELS,
)

batch_size = 2
mstcn_dataloader = DataLoader(
    mstcn_dataset,
    batch_size=batch_size,
    shuffle=True,
    collate_fn=collate_fn,
)

print(f"\nMSTCNDataset created with {len(mstcn_dataset)} samples.")
print(f"DataLoader created with batch size {batch_size}.")

# Test loading a batch:
for batch_features, batch_labels in mstcn_dataloader:
    print(f"\nSample Batch Features shape: {batch_features.shape}")
    print(f"Sample Batch Labels shape: {batch_labels.shape}")
    print("Sample Batch Labels (first sample, first 20 frames):",
batch_labels[0, :20])
    break

print("Data loading utilities successfully prepared with binary frame-
level labels.")

Found 16 raw labels: ['chewing', 'drink', 'eat it', 'food in hand at
table', 'move hand away from mouth', 'move hand towards mouth', 'no
action', 'other', 'pick food from utensil with both hands', 'pick food
from utensil with one hand', 'pick food from utensil with tool in one
hand', 'pick food from utensil with tools in both hands', 'pick up a
cup/glass', 'pick up tools with both hands', 'put one tool back', 'put
the cup/glass back']
Binary label_to_id: {'NON_INTAKE': 0, 'INTAKE': 1}
Binary id_to_label: {0: 'NON_INTAKE', 1: 'INTAKE'}

MSTCNDataset created with 135 samples.
DataLoader created with batch size 2.

```

```
Sample Batch Features shape: torch.Size([2, 24, 9949])
Sample Batch Labels shape: torch.Size([2, 9949])
Sample Batch Labels (first sample, first 20 frames): tensor([0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])
Data loading utilities successfully prepared with binary frame-level
labels.
```

## Intake & NON-Take Distribution in Frame-Level

```
frame_counts = Counter()

label_files = sorted((labels_output_dir).glob("*.txt"))
for label_path in label_files:
    with open(label_path, "r") as f:
        for line in f:
            raw_label = line.strip()
            binary_name = "INTAKE" if raw_label in INTAKE_LABELS else
"NON_INTAKE"
            frame_counts[binary_name] += 1

total_frames = sum(frame_counts.values())

print("Frame counts (all files):")
for cls in ["INTAKE", "NON_INTAKE"]:
    count = frame_counts[cls]
    pct = 100.0 * count / total_frames if total_frames > 0 else 0.0
    print(f"{cls}: {count} frames ({pct:.2f}%)")

print(f"\nTotal frames across all files: {total_frames}")

Frame counts (all files):
INTAKE: 34287 frames (4.83%)
NON_INTAKE: 675540 frames (95.17%)

Total frames across all files: 709827

per_file_counts = {}

for label_path in label_files:
    file_counter = Counter()
    with open(label_path, "r") as f:
        for line in f:
            raw_label = line.strip()
            binary_name = "INTAKE" if raw_label in INTAKE_LABELS else
"NON_INTAKE"
            file_counter[binary_name] += 1
    per_file_counts[label_path.stem] = file_counter

for fname, cnt in sorted(per_file_counts.items()):
```

```
total = sum(cnt.values())
intake = cnt.get("INTAKE", 0)
non_intake = cnt.get("NON_INTAKE", 0)
print(
    f"{fname}: INTAKE={intake} ({intake/total:.2%}), "
    f"NON_INTAKE={non_intake} ({non_intake/total:.2%}),
total={total}"
)

20210518_230219: INTAKE=206 (5.62%), NON_INTAKE=3462 (94.38%),
total=3668
20210523_202300: INTAKE=363 (6.72%), NON_INTAKE=5042 (93.28%),
total=5405
20210529_150552: INTAKE=1055 (14.07%), NON_INTAKE=6443 (85.93%),
total=7498
20210529_153708: INTAKE=0 (0.00%), NON_INTAKE=1395 (100.00%),
total=1395
20210530_153343: INTAKE=490 (9.65%), NON_INTAKE=4590 (90.35%),
total=5080
20210531_150448: INTAKE=763 (13.94%), NON_INTAKE=4711 (86.06%),
total=5474
20210603_130948: INTAKE=136 (1.79%), NON_INTAKE=7480 (98.21%),
total=7616
20210605_155355: INTAKE=285 (7.76%), NON_INTAKE=3389 (92.24%),
total=3674
20210606_154234: INTAKE=440 (9.00%), NON_INTAKE=4450 (91.00%),
total=4890
20210607_152559: INTAKE=283 (11.21%), NON_INTAKE=2241 (88.79%),
total=2524
20210607_154140: INTAKE=0 (0.00%), NON_INTAKE=2534 (100.00%),
total=2534
20210608_163819: INTAKE=137 (1.90%), NON_INTAKE=7056 (98.10%),
total=7193
20210609_154241: INTAKE=513 (8.48%), NON_INTAKE=5539 (91.52%),
total=6052
20210609_215756: INTAKE=142 (2.42%), NON_INTAKE=5731 (97.58%),
total=5873
20210609_220502: INTAKE=112 (9.82%), NON_INTAKE=1028 (90.18%),
total=1140
20210609_221133: INTAKE=0 (0.00%), NON_INTAKE=1230 (100.00%),
total=1230
20210610_144155: INTAKE=287 (4.83%), NON_INTAKE=5653 (95.17%),
total=5940
20210610_230110: INTAKE=158 (1.86%), NON_INTAKE=8344 (98.14%),
total=8502
20210616_163847: INTAKE=245 (5.65%), NON_INTAKE=4092 (94.35%),
total=4337
20210616_170258: INTAKE=0 (0.00%), NON_INTAKE=3144 (100.00%),
total=3144
20210617_165755: INTAKE=253 (3.88%), NON_INTAKE=6276 (96.12%),
```

total=6529  
20210617\_224119: INTAKE=10 (1.98%), NON\_INTAKE=494 (98.02%), total=504  
20210620\_150313: INTAKE=117 (1.38%), NON\_INTAKE=8376 (98.62%),  
total=8493  
20210620\_230306: INTAKE=137 (6.29%), NON\_INTAKE=2041 (93.71%),  
total=2178  
20210620\_230707: INTAKE=175 (4.28%), NON\_INTAKE=3914 (95.72%),  
total=4089  
20210621\_125816: INTAKE=170 (3.63%), NON\_INTAKE=4518 (96.37%),  
total=4688  
20210627\_142532: INTAKE=264 (3.54%), NON\_INTAKE=7202 (96.46%),  
total=7466  
20210714\_125156: INTAKE=343 (5.32%), NON\_INTAKE=6106 (94.68%),  
total=6449  
20210716\_151219: INTAKE=136 (1.87%), NON\_INTAKE=7151 (98.13%),  
total=7287  
20210719\_135251: INTAKE=222 (3.43%), NON\_INTAKE=6245 (96.57%),  
total=6467  
20210728\_144106: INTAKE=220 (2.21%), NON\_INTAKE=9755 (97.79%),  
total=9975  
20210804\_144735: INTAKE=135 (1.73%), NON\_INTAKE=7664 (98.27%),  
total=7799  
20210804\_145925: INTAKE=258 (5.79%), NON\_INTAKE=4199 (94.21%),  
total=4457  
20210810\_163907: INTAKE=441 (12.15%), NON\_INTAKE=3189 (87.85%),  
total=3630  
20210811\_130812: INTAKE=250 (2.51%), NON\_INTAKE=9699 (97.49%),  
total=9949  
20210811\_132135: INTAKE=491 (6.09%), NON\_INTAKE=7574 (93.91%),  
total=8065  
20210813\_195607: INTAKE=77 (2.31%), NON\_INTAKE=3253 (97.69%),  
total=3330  
20210814\_210449: INTAKE=103 (1.64%), NON\_INTAKE=6175 (98.36%),  
total=6278  
20210816\_191136: INTAKE=377 (4.06%), NON\_INTAKE=8903 (95.94%),  
total=9280  
20210923\_155848: INTAKE=198 (2.17%), NON\_INTAKE=8930 (97.83%),  
total=9128  
20210923\_161107: INTAKE=458 (5.07%), NON\_INTAKE=8582 (94.93%),  
total=9040  
20210923\_163035: INTAKE=168 (2.83%), NON\_INTAKE=5769 (97.17%),  
total=5937  
20210923\_163801: INTAKE=257 (4.10%), NON\_INTAKE=6014 (95.90%),  
total=6271  
20211001\_152746: INTAKE=438 (3.53%), NON\_INTAKE=11975 (96.47%),  
total=12413  
20211001\_154310: INTAKE=417 (3.24%), NON\_INTAKE=12469 (96.76%),  
total=12886  
20211029\_144948: INTAKE=391 (6.06%), NON\_INTAKE=6065 (93.94%),

total=6456  
20211029\_150526: INTAKE=84 (3.14%) , NON\_INTAKE=2590 (96.86%) ,  
total=2674  
20220809\_125236: INTAKE=382 (3.05%) , NON\_INTAKE=12158 (96.95%) ,  
total=12540  
20220809\_142839: INTAKE=503 (6.88%) , NON\_INTAKE=6812 (93.12%) ,  
total=7315  
20220809\_144052: INTAKE=461 (6.38%) , NON\_INTAKE=6765 (93.62%) ,  
total=7226  
20220811\_113712: INTAKE=260 (5.12%) , NON\_INTAKE=4818 (94.88%) ,  
total=5078  
20220811\_114634: INTAKE=160 (3.22%) , NON\_INTAKE=4802 (96.78%) ,  
total=4962  
20220811\_121550: INTAKE=294 (6.49%) , NON\_INTAKE=4238 (93.51%) ,  
total=4532  
20220811\_122341: INTAKE=467 (4.85%) , NON\_INTAKE=9156 (95.15%) ,  
total=9623  
20220811\_124226: INTAKE=229 (1.47%) , NON\_INTAKE=15348 (98.53%) ,  
total=15577  
20220812\_112034: INTAKE=448 (15.19%) , NON\_INTAKE=2502 (84.81%) ,  
total=2950  
20220812\_113020: INTAKE=685 (10.08%) , NON\_INTAKE=6110 (89.92%) ,  
total=6795  
20220812\_115050: INTAKE=218 (3.21%) , NON\_INTAKE=6575 (96.79%) ,  
total=6793  
20220812\_120241: INTAKE=180 (2.31%) , NON\_INTAKE=7613 (97.69%) ,  
total=7793  
20220812\_121833: INTAKE=152 (3.33%) , NON\_INTAKE=4408 (96.67%) ,  
total=4560  
20220812\_122535: INTAKE=237 (4.34%) , NON\_INTAKE=5219 (95.66%) ,  
total=5456  
20220812\_123723: INTAKE=111 (4.55%) , NON\_INTAKE=2326 (95.45%) ,  
total=2437  
20220812\_124345: INTAKE=265 (5.67%) , NON\_INTAKE=4409 (94.33%) ,  
total=4674  
20220812\_125132: INTAKE=339 (4.27%) , NON\_INTAKE=7604 (95.73%) ,  
total=7943  
20220812\_141747: INTAKE=643 (5.46%) , NON\_INTAKE=11134 (94.54%) ,  
total=11777  
20221013\_124150: INTAKE=247 (3.34%) , NON\_INTAKE=7153 (96.66%) ,  
total=7400  
20221013\_125118: INTAKE=310 (3.34%) , NON\_INTAKE=8975 (96.66%) ,  
total=9285  
20221013\_131254: INTAKE=149 (1.63%) , NON\_INTAKE=8972 (98.37%) ,  
total=9121  
20221013\_132443: INTAKE=243 (2.62%) , NON\_INTAKE=9023 (97.38%) ,  
total=9266  
20221013\_134216: INTAKE=214 (7.39%) , NON\_INTAKE=2681 (92.61%) ,  
total=2895

20221013\_135010: INTAKE=153 (4.85%), NON\_INTAKE=3000 (95.15%),  
total=3153  
20221013\_142709: INTAKE=193 (6.39%), NON\_INTAKE=2826 (93.61%),  
total=3019  
20221013\_143133: INTAKE=55 (6.01%), NON\_INTAKE=860 (93.99%), total=915  
20221014\_113751: INTAKE=365 (4.24%), NON\_INTAKE=8241 (95.76%),  
total=8606  
20221014\_114952: INTAKE=120 (2.93%), NON\_INTAKE=3978 (97.07%),  
total=4098  
20221014\_120757: INTAKE=258 (2.84%), NON\_INTAKE=8841 (97.16%),  
total=9099  
20221014\_122331: INTAKE=299 (2.62%), NON\_INTAKE=11133 (97.38%),  
total=11432  
20221014\_124303: INTAKE=984 (10.84%), NON\_INTAKE=8097 (89.16%),  
total=9081  
20221014\_125502: INTAKE=1773 (15.76%), NON\_INTAKE=9475 (84.24%),  
total=11248  
20221014\_140942: INTAKE=509 (10.13%), NON\_INTAKE=4518 (89.87%),  
total=5027  
20221014\_141711: INTAKE=556 (5.63%), NON\_INTAKE=9328 (94.37%),  
total=9884  
20221014\_143649: INTAKE=209 (5.62%), NON\_INTAKE=3510 (94.38%),  
total=3719  
20221014\_144233: INTAKE=471 (6.45%), NON\_INTAKE=6830 (93.55%),  
total=7301  
20230109\_114530: INTAKE=178 (4.68%), NON\_INTAKE=3628 (95.32%),  
total=3806  
20230109\_115040: INTAKE=166 (4.08%), NON\_INTAKE=3900 (95.92%),  
total=4066  
20230109\_115716: INTAKE=166 (4.19%), NON\_INTAKE=3796 (95.81%),  
total=3962  
20230109\_120255: INTAKE=86 (3.12%), NON\_INTAKE=2668 (96.88%),  
total=2754  
20230109\_123622: INTAKE=152 (5.89%), NON\_INTAKE=2430 (94.11%),  
total=2582  
20230109\_124022: INTAKE=273 (6.28%), NON\_INTAKE=4073 (93.72%),  
total=4346  
20230109\_124729: INTAKE=180 (6.01%), NON\_INTAKE=2816 (93.99%),  
total=2996  
20230109\_125206: INTAKE=97 (6.69%), NON\_INTAKE=1354 (93.31%),  
total=1451  
20230109\_132632: INTAKE=137 (3.91%), NON\_INTAKE=3370 (96.09%),  
total=3507  
20230109\_133112: INTAKE=135 (4.18%), NON\_INTAKE=3095 (95.82%),  
total=3230  
20230109\_133609: INTAKE=139 (3.83%), NON\_INTAKE=3490 (96.17%),  
total=3629  
20230109\_134123: INTAKE=167 (4.36%), NON\_INTAKE=3663 (95.64%),  
total=3830

20230110\_113650: INTAKE=146 (9.44%), NON\_INTAKE=1400 (90.56%),  
total=1546  
20230110\_113938: INTAKE=154 (7.95%), NON\_INTAKE=1783 (92.05%),  
total=1937  
20230110\_114318: INTAKE=179 (10.25%), NON\_INTAKE=1568 (89.75%),  
total=1747  
20230110\_114618: INTAKE=136 (5.91%), NON\_INTAKE=2165 (94.09%),  
total=2301  
20230110\_122933: INTAKE=202 (7.56%), NON\_INTAKE=2470 (92.44%),  
total=2672  
20230110\_123332: INTAKE=137 (6.28%), NON\_INTAKE=2045 (93.72%),  
total=2182  
20230110\_123733: INTAKE=164 (6.03%), NON\_INTAKE=2555 (93.97%),  
total=2719  
20230110\_124143: INTAKE=123 (4.89%), NON\_INTAKE=2393 (95.11%),  
total=2516  
20230110\_133309: INTAKE=245 (5.02%), NON\_INTAKE=4640 (94.98%),  
total=4885  
20230110\_133944: INTAKE=229 (3.66%), NON\_INTAKE=6027 (96.34%),  
total=6256  
20230110\_134822: INTAKE=193 (4.57%), NON\_INTAKE=4026 (95.43%),  
total=4219  
20230110\_135409: INTAKE=228 (4.30%), NON\_INTAKE=5074 (95.70%),  
total=5302  
20230111\_114057: INTAKE=101 (4.06%), NON\_INTAKE=2385 (95.94%),  
total=2486  
20230111\_114640: INTAKE=126 (5.46%), NON\_INTAKE=2181 (94.54%),  
total=2307  
20230111\_115239: INTAKE=129 (4.58%), NON\_INTAKE=2688 (95.42%),  
total=2817  
20230111\_115909: INTAKE=248 (10.90%), NON\_INTAKE=2028 (89.10%),  
total=2276  
20230124\_121223: INTAKE=140 (4.47%), NON\_INTAKE=2995 (95.53%),  
total=3135  
20230124\_121646: INTAKE=117 (4.19%), NON\_INTAKE=2676 (95.81%),  
total=2793  
20230124\_122123: INTAKE=256 (6.10%), NON\_INTAKE=3940 (93.90%),  
total=4196  
20230124\_122705: INTAKE=95 (3.37%), NON\_INTAKE=2728 (96.63%),  
total=2823  
20230124\_130607: INTAKE=169 (4.84%), NON\_INTAKE=3322 (95.16%),  
total=3491  
20230124\_131109: INTAKE=156 (3.31%), NON\_INTAKE=4554 (96.69%),  
total=4710  
20230124\_131740: INTAKE=184 (4.48%), NON\_INTAKE=3926 (95.52%),  
total=4110  
20230124\_132312: INTAKE=141 (3.54%), NON\_INTAKE=3840 (96.46%),  
total=3981  
20230124\_135539: INTAKE=167 (5.84%), NON\_INTAKE=2695 (94.16%),

```

total=2862
20230124_135936: INTAKE=155 (4.76%) , NON_INTAKE=3101 (95.24%) ,
total=3256
20230124_140432: INTAKE=220 (4.00%) , NON_INTAKE=5281 (96.00%) ,
total=5501
20230124_141139: INTAKE=132 (4.29%) , NON_INTAKE=2946 (95.71%) ,
total=3078
20230125_140902: INTAKE=138 (3.38%) , NON_INTAKE=3947 (96.62%) ,
total=4085
20230125_141519: INTAKE=114 (2.99%) , NON_INTAKE=3704 (97.01%) ,
total=3818
20230125_142103: INTAKE=128 (3.44%) , NON_INTAKE=3593 (96.56%) ,
total=3721
20230125_142618: INTAKE=99 (3.76%) , NON_INTAKE=2534 (96.24%) ,
total=2633
20230127_120259: INTAKE=163 (6.32%) , NON_INTAKE=2417 (93.68%) ,
total=2580
20230127_120640: INTAKE=217 (6.10%) , NON_INTAKE=3343 (93.90%) ,
total=3560
20230127_121212: INTAKE=347 (7.31%) , NON_INTAKE=4401 (92.69%) ,
total=4748
20230127_121834: INTAKE=175 (5.06%) , NON_INTAKE=3282 (94.94%) ,
total=3457
20230131_130800: INTAKE=224 (6.18%) , NON_INTAKE=3401 (93.82%) ,
total=3625
20230131_131313: INTAKE=229 (4.96%) , NON_INTAKE=4387 (95.04%) ,
total=4616
20230131_131949: INTAKE=299 (4.88%) , NON_INTAKE=5823 (95.12%) ,
total=6122
20230131_132739: INTAKE=259 (3.63%) , NON_INTAKE=6873 (96.37%) ,
total=7132

```

### Summary Highlights:

- Out of 135 videos, 4 sessions had 0% INTAKE (no "eat it" frames).
- We kept these files because they still provide valid NON-INTAKE behavior (talking, hand movements, background motion) that the model must learn to ignore.
- Removing them would inflate the apparent intake rate and bias the dataset toward "eating" clips, making real-world performance less realistic.
- Including true "no-intake" sessions also gives a harder, more honest evaluation, since the model must correctly return zero bites for some videos.

```

features_output_dir = Path(features_output_dir)
labels_output_dir = Path(labels_output_dir)

all_feature_files = sorted(features_output_dir.glob("*.npy"))
all_filenames = [f.stem for f in all_feature_files]

per_file_stats = []

```

```

for fname in all_filenames:
    label_path = labels_output_dir / f"{fname}.txt"
    cnt = Counter()
    with open(label_path, "r") as f:
        for line in f:
            raw_label = line.strip()
            cls_name = "INTAKE" if raw_label in INTAKE_LABELS else
"NON_INTAKE"
            cnt[cls_name] += 1
    intake = cnt["INTAKE"]
    non_intake = cnt["NON_INTAKE"]
    total = intake + non_intake
    ratio = intake / total if total > 0 else 0.0
    per_file_stats.append(
        {"filename": fname, "intake": intake, "non_intake":
non_intake,
         "total": total, "ratio": ratio}
    )

# Sanity:
global_intake = sum(s["intake"] for s in per_file_stats)
global_total = sum(s["total"] for s in per_file_stats)
print(f"Global INTAKE%: {100 * global_intake / global_total:.2f}%")

Global INTAKE%: 4.83%

```

### Summary Highlights:

- With frame-level labels, about 5% of frames are INTAKE, vs only 0.4% positive windows in the earlier pose-TCN setup.
- This means the class imbalance is still strong but less extreme, giving MS-TCN a richer positive signal to learn from.
- Practically, we now see more “eat it” examples per video, which should help the model better capture the timing and shape of intake events.

### Split: Train, Test, Val

```

ratios = np.array([s["ratio"] for s in per_file_stats])
bins = [0.0, 0.01, 0.03, 0.06, 0.10, 1.0]
ratio_bins = np.digitize(ratios, bins)

filenames = np.array([s["filename"] for s in per_file_stats])

# splits ( 70 / 15 / 15):
test_size = 0.15
val_size = 0.15
train_size = 1.0 - test_size - val_size

# Stratified train vs temp (val+test):

```

```

sssl = StratifiedShuffleSplit(
    n_splits=1,
    test_size=(val_size + test_size),
    random_state=42
)
train_idx, temp_idx = next(sssl.split(filenames, ratio_bins))

X_train = filenames[train_idx]
y_train_bins = ratio_bins[train_idx]

X_temp = filenames[temp_idx]
y_temp_bins = ratio_bins[temp_idx]

# Random val vs test from temp:
relative_test_size = test_size / (test_size + val_size)
X_val, X_test, y_val_bins, y_test_bins = train_test_split(
    X_temp,
    y_temp_bins,
    test_size=relative_test_size,
    random_state=43,
    shuffle=True,
    stratify=None,
)

print(f"Train files: {len(X_train)}, Val files: {len(X_val)}, Test files: {len(X_test)}")

# let's check intake% in each split:
def split_intake_ratio(filenames_subset):
    intake = 0
    total = 0
    lookup = {s["filename"]: s for s in per_file_stats}
    for fn in filenames_subset:
        stats = lookup[fn]
        intake += stats["intake"]
        total += stats["total"]
    return 100 * intake / total

print(f"Train INTAKE%: {split_intake_ratio(X_train):.2f}")
print(f"Val INTAKE%: {split_intake_ratio(X_val):.2f}")
print(f"Test INTAKE%: {split_intake_ratio(X_test):.2f}")

Train files: 94, Val files: 20, Test files: 21
Train INTAKE%: 4.78
Val INTAKE%: 5.82
Test INTAKE%: 3.98

```

## MS-TCN Model Training

```
class MSTCNDataset(Dataset):
    def __init__(self, features_dir, labels_dir, label_to_id,
intake_labels, filenames):
        self.features_dir = Path(features_dir)
        self.labels_dir = Path(labels_dir)
        self.label_to_id = label_to_id
        self.intake_labels = intake_labels
        self.filenames = sorted(list(filenames))

    def __len__(self):
        return len(self.filenames)

    def __getitem__(self, idx):
        filename = self.filenames[idx]

        # Load features:
        features_path = self.features_dir / f"{filename}.npy"
        features = np.load(features_path)
        features_tensor = torch.from_numpy(features).float() # (T, F)
        features_tensor = features_tensor.transpose(0, 1) # (F, T)

        # Load frame-level labels:
        labels_path = self.labels_dir / f"{filename}.txt"
        with open(labels_path, "r") as f:
            raw_labels = [line.strip() for line in f]

        # Map to binary classes:
        numerical_labels = [
            self.label_to_id["INTAKE"] if lbl in self.intake_labels
            else self.label_to_id["NON_INTAKE"]
            for lbl in raw_labels
        ]
        labels_tensor = torch.tensor(numerical_labels,
dtype=torch.long) # (T,)

        return features_tensor, labels_tensor

batch_size = 2

train_dataset = MSTCNDataset(features_output_dir, labels_output_dir,
                             label_to_id, INTAKE_LABELS, X_train)
val_dataset   = MSTCNDataset(features_output_dir, labels_output_dir,
                             label_to_id, INTAKE_LABELS, X_val)
test_dataset  = MSTCNDataset(features_output_dir, labels_output_dir,
                             label_to_id, INTAKE_LABELS, X_test)

train_loader = DataLoader(train_dataset, batch_size=batch_size,
                         shuffle=True, collate_fn=collate_fn)
val_loader   = DataLoader(val_dataset, batch_size=batch_size,
```

```

        shuffle=False, collate_fn=collate_fn)
test_loader = DataLoader(test_dataset, batch_size=batch_size,
                        shuffle=False, collate_fn=collate_fn)

print(f"Train samples: {len(train_dataset)}")
print(f"Val samples: {len(val_dataset)}")
print(f"Test samples: {len(test_dataset)}")

Train samples: 94
Val samples: 20
Test samples: 21

import collections

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

num_classes = len(label_to_id)

label_counts = collections.defaultdict(int)

for fname in X_train:
    labels_path = labels_output_dir / f"{fname}.txt"
    with open(labels_path, "r") as f:
        for line in f:
            raw_label = line.strip()
            cls_name = "INTAKE" if raw_label in INTAKE_LABELS else
"NON_INTAKE"
            cls_id = label_to_id[cls_name]
            label_counts[cls_id] += 1

class_frequencies = torch.zeros(num_classes, dtype=torch.float32)
for label_id, count in label_counts.items():
    class_frequencies[label_id] = count

total_samples_in_labels = class_frequencies.sum()
class_weights = total_samples_in_labels / (class_frequencies * num_classes)
class_weights[class_frequencies == 0] = 0.0
class_weights = class_weights.to(device)

print("Train class frequencies:", class_frequencies)
print("Train class weights:", class_weights)

Train class frequencies: tensor([481052., 24169.])
Train class weights: tensor([ 0.5251, 10.4518], device='cuda:0')

class ConvBlock(nn.Module):
    def __init__(self, in_channels, out_channels, kernel_size,
dilation=1):
        super(ConvBlock, self).__init__()
        padding = self._calculate_padding(kernel_size, dilation)

```

```

        self.conv = nn.Conv1d(
            in_channels,
            out_channels,
            kernel_size,
            padding=padding,
            dilation=dilation
        )
        self.bn = nn.BatchNorm1d(out_channels)
        self.relu = nn.ReLU(inplace=True)

    def _calculate_padding(self, kernel_size, dilation):
        return dilation * (kernel_size - 1) // 2

    def forward(self, x):
        x = self.conv(x)
        x = self.bn(x)
        x = self.relu(x)
        return x

class DilatedResidualBlock(nn.Module):
    def __init__(self, in_channels, out_channels, kernel_size,
dilation):
        super(DilatedResidualBlock, self).__init__()
        self.conv1 = ConvBlock(in_channels, out_channels, kernel_size,
dilation)
        self.conv2 = ConvBlock(out_channels, out_channels,
kernel_size, dilation)
        self.downsample = None
        if in_channels != out_channels:
            self.downsample = nn.Conv1d(in_channels, out_channels,
kernel_size=1)

    def forward(self, x):
        residual = x
        if self.downsample:
            residual = self.downsample(residual)

        out = self.conv1(x)
        out = self.conv2(out)
        return out + residual

class SingleStageTCN(nn.Module):
    def __init__(self, num_layers, in_channels, out_channels,
kernel_size):
        super(SingleStageTCN, self).__init__()
        self.layers = nn.ModuleList()
        for i in range(num_layers):
            # Increase dilation factor for each layer
            dilation = 2 ** i
            self.layers.append(

```

```

        DilatedResidualBlock(in_channels if i == 0 else
out_channels, out_channels, kernel_size, dilation)
    )

def forward(self, x):
    for layer in self.layers:
        x = layer(x)
    return x

# Here's the main MS_TCN class:
class MS_TCN(nn.Module):
    def __init__(self, num_stages, num_layers_per_stage, in_features,
num_classes, kernel_size=3):
        super(MS_TCN, self).__init__()
        self.num_stages = num_stages
        self.stages = nn.ModuleList()
        # The number of channels output by each stage:
        stage_out_channels = 64

        # Initial convolution to project input features to
stage_out_channels:
        self.initial_conv = nn.Conv1d(in_features, stage_out_channels,
kernel_size=1)

        for s in range(num_stages):
            self.stages.append(
                SingleStageTCN(num_layers_per_stage,
stage_out_channels, stage_out_channels, kernel_size)
            )

        # Output layer for frame-level classification: 1x1
convolution:
        self.output_conv = nn.Conv1d(stage_out_channels, num_classes,
kernel_size=1)

    def forward(self, x):
        # (batch_size, in_features, sequence_length)
        # Initial convolution:
        x = self.initial_conv(x)

        for stage in self.stages:
            x = stage(x)

        # Apply 1x1 convolution for frame-level classification
        # Output = (batch_size, num_classes, sequence_length)
        out = self.output_conv(x)
        return out

# Re-instantiate the model with the new architecture:
in_features = mstcn_dataset[0][0].shape[0]

```

```

num_classes = len(label_to_id) # Number of unique labels

# Define model hyperparameters:
num_stages = 2
num_layers_per_stage = 5

model = MS_TCN(num_stages, num_layers_per_stage, in_features,
num_classes).to(device)
criterion = nn.CrossEntropyLoss(ignore_index=-100,
weight=class_weights)
optimizer = optim.Adam(model.parameters(), lr=0.001)

def train_model_frame_level(
    model,
    train_loader,
    val_loader,
    criterion,
    optimizer,
    num_epochs=50,
    patience=5,
    min_delta=0.0,
    checkpoint_path=None,
):
    """
    Frame-level training with early stopping on validation loss.
    - patience: number of epochs with no improvement before stopping
    - min_delta: minimum change in val loss to qualify as improvement
    """
    print(f"\nStarting frame-level training for up to {num_epochs} epochs...")
    best_val_loss = math.inf
    best_model_state = copy.deepcopy(model.state_dict())
    epochs_no_improve = 0

    for epoch in range(num_epochs):
        # ----- TRAIN -----
        model.train()
        epoch_loss = 0.0
        start_time = time.time()

        for features, labels in train_loader:
            features = features.to(device)
            labels = labels.to(device)

            optimizer.zero_grad()

            outputs = model(features)                      # (B, C, T)
            outputs_reshaped = outputs.permute(0, 2, 1)    # (B, T, C)
            outputs_flat = outputs_reshaped.reshape(-1, num_classes)

```

```

        labels_flat = labels.reshape(-1)

        loss = criterion(outputs_flat, labels_flat)
        loss.backward()
        optimizer.step()

        epoch_loss += loss.item()

    avg_train_loss = epoch_loss / len(train_loader)

    # VALIDATION:
    model.eval()
    val_loss = 0.0
    with torch.no_grad():
        for features, labels in val_loader:
            features = features.to(device)
            labels = labels.to(device)

            outputs = model(features)
            outputs_reshaped = outputs.permute(0, 2, 1)
            outputs_flat = outputs_reshaped.reshape(-1,
num_classes)
            labels_flat = labels.reshape(-1)

            loss = criterion(outputs_flat, labels_flat)
            val_loss += loss.item()

    avg_val_loss = val_loss / len(val_loader)
    elapsed = time.time() - start_time

    print(
        f"Epoch {epoch+1}/{num_epochs} | "
        f"Train Loss: {avg_train_loss:.4f} | "
        f"Val Loss: {avg_val_loss:.4f} | "
        f"Time: {elapsed:.2f}s"
    )

    # EARLY STOPPING CHECK:
    if avg_val_loss + min_delta < best_val_loss:
        best_val_loss = avg_val_loss
        best_model_state = copy.deepcopy(model.state_dict())
        epochs_no_improve = 0

        # checkpoint:
        if checkpoint_path is not None:
            torch.save(best_model_state, checkpoint_path)
    else:
        epochs_no_improve += 1
        print(f" No improvement for {epochs_no_improve}
epoch(s).")

```

```

        if epochs_no_improve >= patience:
            print(
                f"\nEarly stopping triggered after {epoch+1} epochs "
                f"(best val loss = {best_val_loss:.4f})."
            )
            break

# Restore best model weights
model.load_state_dict(best_model_state)
print("Frame-level training complete. Best model restored.")

return model

checkpoint_path = "best_mstcn_frame_level_1.pth"

model = MS_TCN(num_stages, num_layers_per_stage, in_features,
num_classes).to(device)
criterion = nn.CrossEntropyLoss(ignore_index=-100,
weight=class_weights)
optimizer = optim.Adam(model.parameters(), lr=0.001)

model = train_model_frame_level(
    model,
    train_loader,
    val_loader,
    criterion,
    optimizer,
    num_epochs=50,
    patience=5,
    min_delta=0.0,
    checkpoint_path=checkpoint_path,
)

```

```

Starting frame-level training for up to 50 epochs...
Epoch 1/50 | Train Loss: 0.5803 | Val Loss: 0.5910 | Time: 89.07s
Epoch 2/50 | Train Loss: 0.4007 | Val Loss: 0.7098 | Time: 2.01s
    No improvement for 1 epoch(s).
Epoch 3/50 | Train Loss: 0.3604 | Val Loss: 0.7582 | Time: 1.79s
    No improvement for 2 epoch(s).
Epoch 4/50 | Train Loss: 0.3046 | Val Loss: 0.4631 | Time: 1.66s
Epoch 5/50 | Train Loss: 0.3159 | Val Loss: 0.6278 | Time: 1.68s
    No improvement for 1 epoch(s).
Epoch 6/50 | Train Loss: 0.2893 | Val Loss: 0.3101 | Time: 1.64s
Epoch 7/50 | Train Loss: 0.3182 | Val Loss: 0.3566 | Time: 1.57s
    No improvement for 1 epoch(s).
Epoch 8/50 | Train Loss: 0.2996 | Val Loss: 0.6153 | Time: 1.57s
    No improvement for 2 epoch(s).

```

```

Epoch 9/50 | Train Loss: 0.2752 | Val Loss: 0.3841 | Time: 1.60s
    No improvement for 3 epoch(s).
Epoch 10/50 | Train Loss: 0.2521 | Val Loss: 0.3384 | Time: 1.62s
    No improvement for 4 epoch(s).
Epoch 11/50 | Train Loss: 0.2721 | Val Loss: 0.4506 | Time: 1.66s
    No improvement for 5 epoch(s).

Early stopping triggered after 11 epochs (best val loss = 0.3101).
Frame-level training complete. Best model restored.

def train_model_frame_level(
    model,
    train_loader,
    val_loader,
    criterion,
    optimizer,
    num_epochs=50,
    patience=5,
    min_delta=0.0,
    checkpoint_path=None,
):
    """
    Frame-level training with early stopping on validation loss.

    Returns:
        - model (restored to best val loss)
        - history = {"train_loss": [...], "val_loss": [...]}

    """
    print(f"\nStarting frame-level training for up to {num_epochs} epochs...")
    best_val_loss = math.inf
    best_model_state = copy.deepcopy(model.state_dict())
    epochs_no_improve = 0

    train_losses = []
    val_losses = []

    for epoch in range(num_epochs):
        # Train:
        model.train()
        epoch_loss = 0.0
        start_time = time.time()

        for features, labels in train_loader:
            features = features.to(device)
            labels = labels.to(device)

            optimizer.zero_grad()

            outputs = model(features)                      # (B, C, T)
            outputs_reshaped = outputs.permute(0, 2, 1)     # (B, T, C)

```

```

outputs_flat = outputs_reshaped.reshape(-1, num_classes)
labels_flat = labels.reshape(-1)

loss = criterion(outputs_flat, labels_flat)
loss.backward()
optimizer.step()

epoch_loss += loss.item()

avg_train_loss = epoch_loss / len(train_loader)
train_losses.append(avg_train_loss)

# VALIDATION:
model.eval()
val_loss = 0.0
with torch.no_grad():
    for features, labels in val_loader:
        features = features.to(device)
        labels = labels.to(device)

        outputs = model(features)
        outputs_reshaped = outputs.permute(0, 2, 1)
        outputs_flat = outputs_reshaped.reshape(-1,
num_classes)
        labels_flat = labels.reshape(-1)

        loss = criterion(outputs_flat, labels_flat)
        val_loss += loss.item()

    avg_val_loss = val_loss / len(val_loader)
    val_losses.append(avg_val_loss)

elapsed = time.time() - start_time

print(
    f"Epoch {epoch+1}/{num_epochs} | "
    f"Train Loss: {avg_train_loss:.4f} | "
    f"Val Loss: {avg_val_loss:.4f} | "
    f"Time: {elapsed:.2f}s"
)

# EARLY STOPPING CHECK:
if avg_val_loss + min_delta < best_val_loss:
    best_val_loss = avg_val_loss
    best_model_state = copy.deepcopy(model.state_dict())
    epochs_no_improve = 0

    if checkpoint_path is not None:
        torch.save(best_model_state, checkpoint_path)
else:

```

```

        epochs_no_improve += 1
        print(f" No improvement for {epochs_no_improve} "
epoch(s).")

        if epochs_no_improve >= patience:
            print(
                f"\nEarly stopping triggered after {epoch+1} "
epochs "
                    f"(best val loss = {best_val_loss:.4f})."
            )
            break

# Restore best model weights:
model.load_state_dict(best_model_state)
print("Frame-level training complete. Best model restored.")

history = {"train_loss": train_losses, "val_loss": val_losses}
return model, history

checkpoint_path = "best_mstcn_frame_level.pth"

model = MS_TCN(num_stages, num_layers_per_stage, in_features,
num_classes).to(device)
criterion = nn.CrossEntropyLoss(ignore_index=-100,
weight=class_weights)
optimizer = optim.Adam(model.parameters(), lr=0.001)

model, history = train_model_frame_level(
    model,
    train_loader,
    val_loader,
    criterion,
    optimizer,
    num_epochs=50,
    patience=5,
    min_delta=0.0,
    checkpoint_path=checkpoint_path,
)

# Plot train vs val loss:
train_losses = history["train_loss"]
val_losses = history["val_loss"]
epochs = range(1, len(train_losses) + 1)

plt.figure(figsize=(8, 5))
plt.plot(epochs, train_losses, marker="o", label="Train Loss")
plt.plot(epochs, val_losses, marker="o", label="Val Loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.title("Train vs Validation Loss (Frame-level, MS-TCN)")

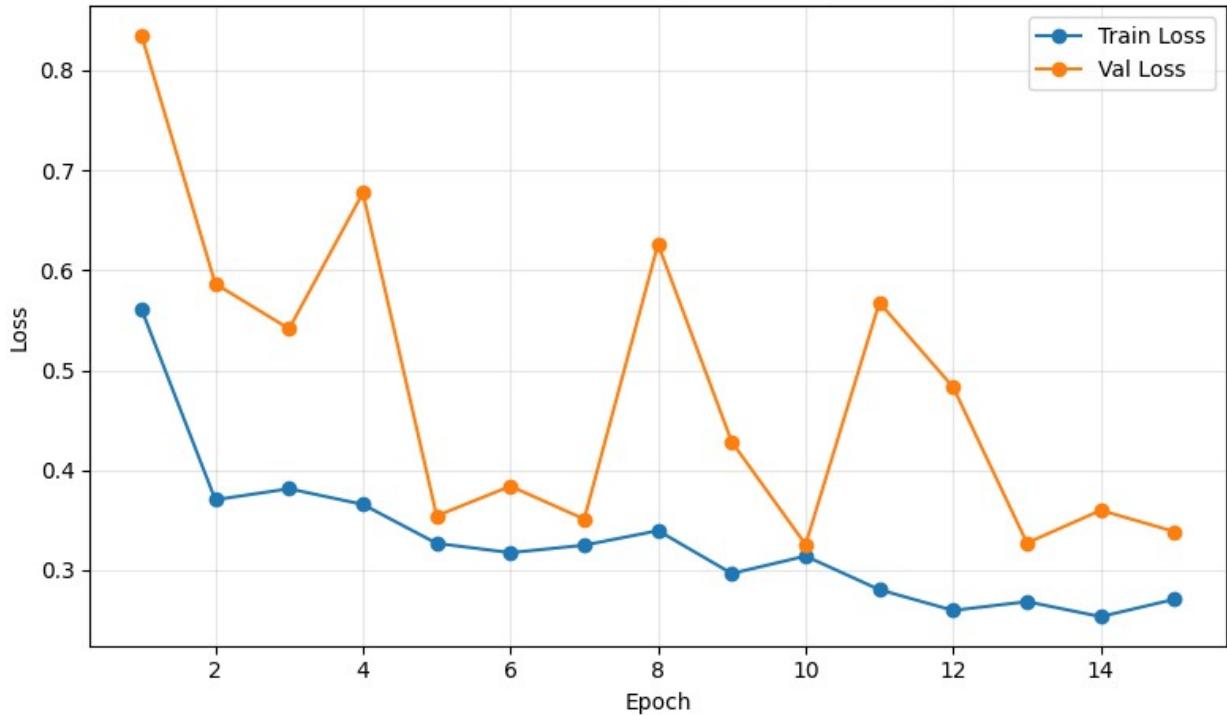
```

```
plt.legend(loc="upper right")
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

Starting frame-level training for up to 50 epochs...
Epoch 1/50 | Train Loss: 0.5611 | Val Loss: 0.8353 | Time: 1.61s
Epoch 2/50 | Train Loss: 0.3704 | Val Loss: 0.5865 | Time: 1.60s
Epoch 3/50 | Train Loss: 0.3814 | Val Loss: 0.5415 | Time: 1.66s
Epoch 4/50 | Train Loss: 0.3657 | Val Loss: 0.6774 | Time: 1.62s
    No improvement for 1 epoch(s).
Epoch 5/50 | Train Loss: 0.3267 | Val Loss: 0.3543 | Time: 1.62s
Epoch 6/50 | Train Loss: 0.3174 | Val Loss: 0.3840 | Time: 1.62s
    No improvement for 1 epoch(s).
Epoch 7/50 | Train Loss: 0.3249 | Val Loss: 0.3510 | Time: 1.61s
Epoch 8/50 | Train Loss: 0.3394 | Val Loss: 0.6263 | Time: 1.66s
    No improvement for 1 epoch(s).
Epoch 9/50 | Train Loss: 0.2966 | Val Loss: 0.4282 | Time: 1.61s
    No improvement for 2 epoch(s).
Epoch 10/50 | Train Loss: 0.3140 | Val Loss: 0.3255 | Time: 1.65s
Epoch 11/50 | Train Loss: 0.2804 | Val Loss: 0.5677 | Time: 1.64s
    No improvement for 1 epoch(s).
Epoch 12/50 | Train Loss: 0.2596 | Val Loss: 0.4830 | Time: 1.60s
    No improvement for 2 epoch(s).
Epoch 13/50 | Train Loss: 0.2684 | Val Loss: 0.3271 | Time: 1.59s
    No improvement for 3 epoch(s).
Epoch 14/50 | Train Loss: 0.2534 | Val Loss: 0.3598 | Time: 1.57s
    No improvement for 4 epoch(s).
Epoch 15/50 | Train Loss: 0.2708 | Val Loss: 0.3385 | Time: 1.60s
    No improvement for 5 epoch(s).

Early stopping triggered after 15 epochs (best val loss = 0.3255).
Frame-level training complete. Best model restored.
```

Train vs Validation Loss (Frame-level, MS-TCN)



```

model.eval()

all_true = []
all_pred = []

with torch.no_grad():
    for features, labels in test_loader:
        features = features.to(device)
        labels = labels.to(device) # (B, T) with -100 for padding

        outputs = model(features) # (B, C, T)
        preds = outputs.argmax(dim=1) # (B, T), class indices

        # Mask out padded positions
        mask = labels != -100

        true_flat = labels[mask].cpu().numpy()
        pred_flat = preds[mask].cpu().numpy()

        all_true.append(true_flat)
        all_pred.append(pred_flat)

y_true = np.concatenate(all_true)
y_pred = np.concatenate(all_pred)

print("True labels shape:", y_true.shape)
print("Pred labels shape:", y_pred.shape)

```

```
True labels shape: (97294, )
Pred labels shape: (97294, )
```

## Model Performance Evaluation

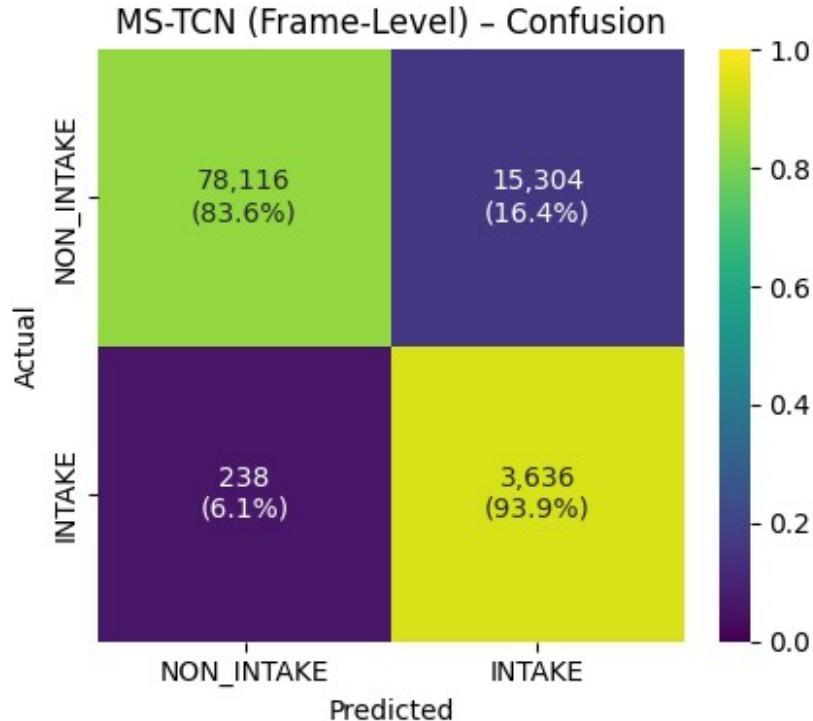
```
# Class names in the correct order
class_names = [id_to_label[0], id_to_label[1]] # ["NON_INTAKE",
"INTAKE"]

# Confusion matrix (float):
cm = confusion_matrix(y_true, y_pred, labels=[0, 1]).astype(float)

# Row-wise percentages (per true class):
row_sums = cm.sum(axis=1, keepdims=True)
row_sums[row_sums == 0] = 1.0
cm_percent = cm / row_sums

# Annotation:
annot = np.empty_like(cm, dtype=object)
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        count = int(cm[i, j])
        pct = cm_percent[i, j] * 100.0
        annot[i, j] = f"{count},\n{pct:.1f}%""

plt.figure(figsize=(4.5, 4))
ax = sns.heatmap(
    cm_percent,
    annot=annot,
    fmt="",
    cmap="viridis",
    vmin=0.0,
    vmax=1.0,
    xticklabels=class_names,
    yticklabels=class_names,
    cbar_kws={"label": ""},
)
ax.set_xlabel("Predicted")
ax.set_ylabel("Actual")
ax.set_title("MS-TCN (Frame-Level) – Confusion")
plt.tight_layout()
plt.show()
```



```

# class_names in the same order as class indices
class_names = [id_to_label[0], id_to_label[1]] # ["NON_INTAKE",
"INTAKE"]

# Overall accuracy:
accuracy = accuracy_score(y_true, y_pred)
print(f"Overall accuracy (frame-level, test set): {accuracy:.4f}\n")

# Detailed per-class metrics:
precision, recall, f1, support = precision_recall_fscore_support(
    y_true, y_pred, labels=[0, 1], average=None
)

print("Per-class KPIs (frame-level, test set):")
for idx, cls_name in enumerate(class_names):
    print(
        f"- {cls_name}: "
        f"Precision = {precision[idx]:.4f}, "
        f"Recall = {recall[idx]:.4f}, "
        f"F1 = {f1[idx]:.4f}, "
        f"Support = {support[idx]}"
    )

# classification_report:
print("\nFull classification report:")
print(classification_report(y_true, y_pred, target_names=class_names,
digits=4))

```

```

Overall accuracy (frame-level, test set): 0.8403

Per-class KPIs (frame-level, test set):
- NON_INTAKE: Precision = 0.9970, Recall = 0.8362, F1 = 0.9095,
Support = 93420
- INTAKE: Precision = 0.1920, Recall = 0.9386, F1 = 0.3188, Support =
3874

Full classification report:
      precision    recall   f1-score   support
NON_INTAKE      0.9970    0.8362    0.9095    93420
INTAKE          0.1920    0.9386    0.3188    3874
accuracy           -         -        0.8403    97294
macro avg       0.5945    0.8874    0.6141    97294
weighted avg    0.9649    0.8403    0.8860    97294

model.eval()

all_true = []
all_scores = [] # probability of INTAKE (class 1)

with torch.no_grad():
    for features, labels in test_loader:
        features = features.to(device)
        labels = labels.to(device) # (B, T) with -100 padding

        outputs = model(features) # (B, C, T)
        probs = torch.softmax(outputs, dim=1) # (B, C, T)

        # Take probability of INTAKE:
        intake_probs = probs[:, 1, :] # (B, T)

        # Mask out padded positions:
        mask = labels != -100

        true_flat = labels[mask].cpu().numpy() # 0 or 1
        score_flat = intake_probs[mask].cpu().numpy() # prob of
INTAKE

        all_true.append(true_flat)
        all_scores.append(score_flat)

y_true = np.concatenate(all_true) # shape (N_frames,)
y_scores = np.concatenate(all_scores) # shape (N_frames,)

print("y_true shape:", y_true.shape)
print("y_scores shape:", y_scores.shape)

```

```

y_true shape: (97294,)
y_scores shape: (97294,)

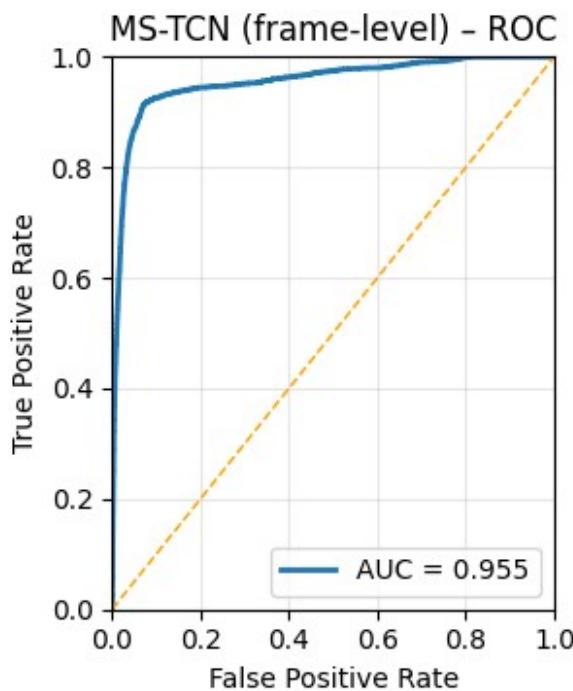
# y_true: 0/1 labels, y_scores: predicted probs for INTAKE:
fpr, tpr, _ = roc_curve(y_true, y_scores, pos_label=1)
roc_auc = roc_auc_score(y_true, y_scores)

plt.figure(figsize=(3.2, 3.8))
plt.plot(fpr, tpr, lw=2, label=f"AUC = {roc_auc:.3f}")
plt.plot([0, 1], [0, 1], linestyle="--", color="orange", lw=1)

plt.xlim(0.0, 1.0)
plt.ylim(0.0, 1.0)
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("MS-TCN (frame-level) – ROC")
plt.grid(alpha=0.3)
plt.legend(loc="lower right")

plt.tight_layout()
plt.show()

```



```

# y_true: 0/1 labels, y_scores: predicted probs for INTAKE
precision, recall, _ = precision_recall_curve(y_true, y_scores,
pos_label=1)
pr_ap = average_precision_score(y_true, y_scores)

plt.figure(figsize=(3.2, 3.8))

```

```

plt.plot(recall, precision, lw=2)

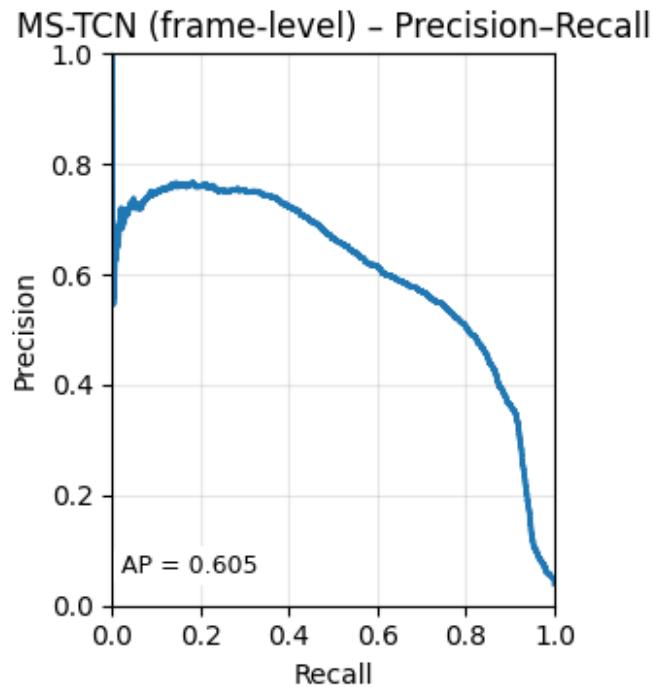
plt.xlim(0.0, 1.0)
plt.ylim(0.0, 1.0)
plt.xlabel("Recall")
plt.ylabel("Precision")
plt.title("MS-TCN (frame-level) – Precision–Recall")
plt.grid(alpha=0.3)

# AP label inside the plot:
plt.text(
    0.02, 0.06,
    f"AP = {pr_ap:.3f}",
    fontsize=9,
    bbox=dict(boxstyle="round,pad=0.3", fc="white", ec="none"),
)

plt.tight_layout()
plt.show()

print(f"PR AUC / AP (INTAKE): {pr_ap:0.4f}")

```



```
PR AUC / AP (INTAKE): 0.6053
```